

# Introduction to Communication Networks and Distributed Systems

Unit 7 – Complex Communication Patterns: Pub/Sub -

Fachgebiet Telekommunikationsnetze

Prof. Dr. A. Zubow

[zubow@tkn.tu-berlin.de](mailto:zubow@tkn.tu-berlin.de)

# Acknowledgements

- We acknowledge the use of slides from: **Prof. Adam Wolisz**
- "Handbook of Research: Ubiquitous Computing Technology for Real Time Enterprises“, Max Mühlhäuser and Iryna Gurevych.
- Eugster, Patrick Th, et al. "The many faces of publish/subscribe." *ACM computing surveys (CSUR)* 35.2 (2003): 114-131.

# Motivation

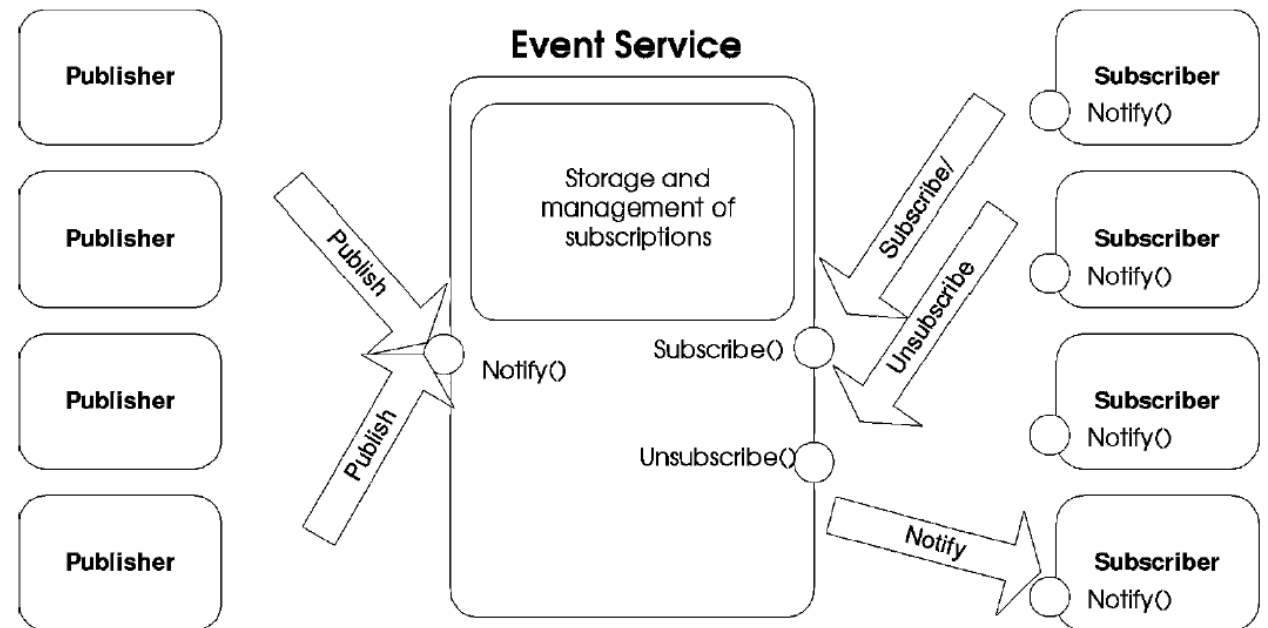
- Traditional client/server communication model (using e.g. RPC, message queue, shared memory, etc.)
  - Synchronous, tightly-coupled request invocations.
  - Very restrictive for distributed applications, especially for Wide Area Network (WAN) and mobile environments.
  - When nodes/links fail, system is affected.
    - **Fault tolerance** must be built in to support this.
- Require a more **flexible** and **de-coupled communication** style that offers anonymous and asynchronous mechanisms.

# What is a publish/subscribe system?

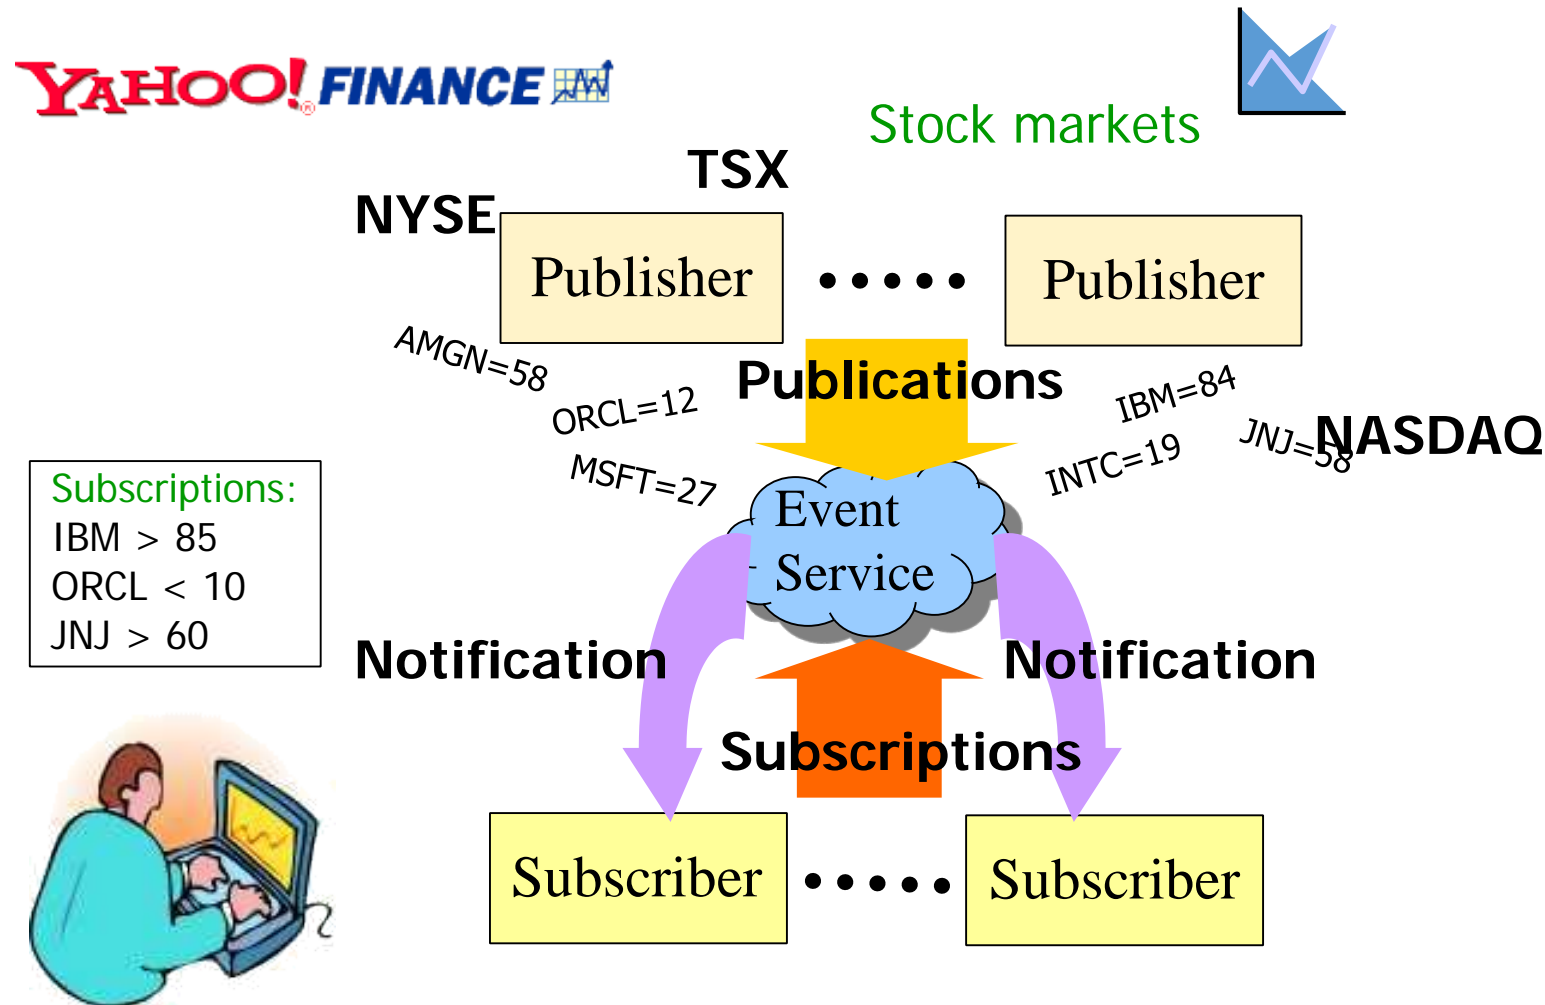
- Distributed pub/sub system is a **communication paradigm** that allows freedom in the distributed system by the **decoupling** of communication entities in terms of **time**, **space** and **synchronization**.
- An **event service system** that is asynchronous, anonymous and loosely-coupled.
- Ability to quickly adapt in a dynamic environment.

# Basic system model for publish/subscribe

- Key components of pub/sub:
  - **Publishers:** Publishers generate event data and publishes them.
  - **Subscribers:** Subscribers submit their subscriptions and process the events received.
  - **Event Service:** It's the mediator/broker that filters and routes events from publishers to interested subscribers.



# Example of pub/sub – Stock Quote Service

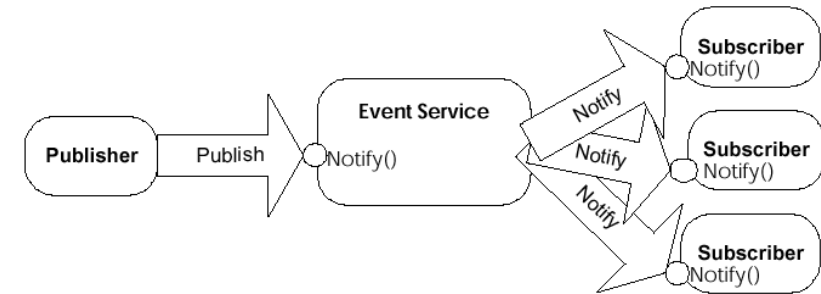


# Towards loosely coupled systems

[Eugster, op. cit.]

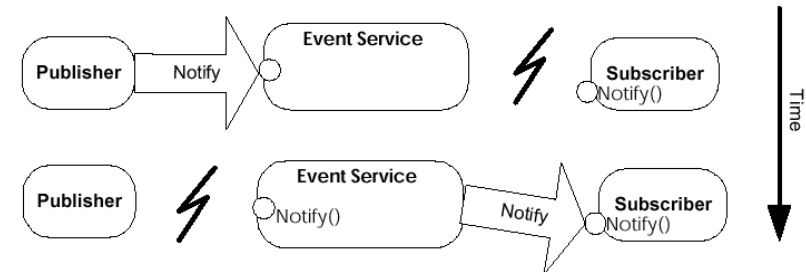
## 1. Space decoupling

- parties don't know each other
- 1-to-many communication possible




## 2. Time decoupling:

- parties not (necessarily) active at same time



## 3. Flow decoupling

- event production & consumption  $\notin$  main control flow (?)
- 1+2+3: coordination & synchronization drastically reduced  $\rightarrow$  **scalability**



# Interaction models

- Interaction models in distributed systems can be classified according to
  - who **initiated** the interaction
  - how the communication partner is **addressed**

	Consumer-initiated („pull“)	Provider-initiated („push“)
Direct Addressing	Request/Reply	Callback
Indirect Addressing	Anonymous Request/Reply	Event-based



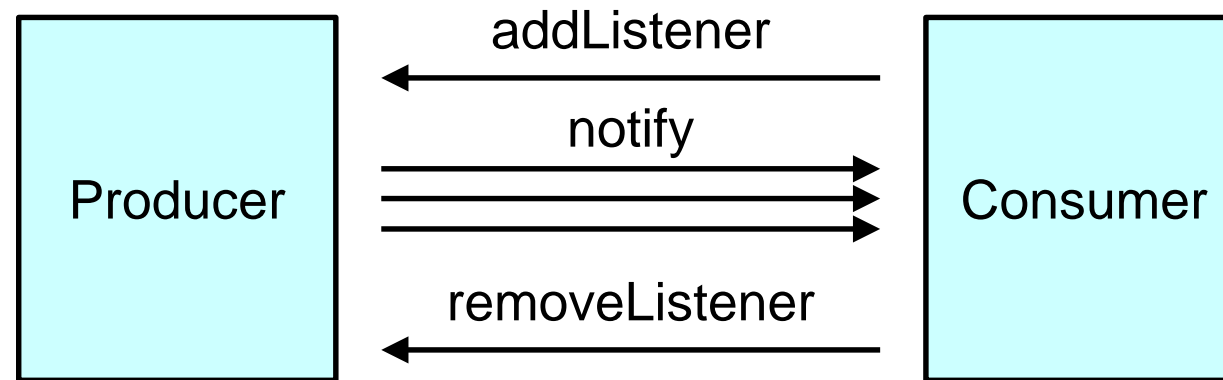
**Today's focus**

- Provider: provides data or functionality
- Anonymous Request/Reply: provider is selected by communication system and not specified directly (e.g., IP Anycast)



# Concepts: Callbacks


- Synchronous (remote) method calls often used to emulate behavior of event-based systems
  - See also: Observer Design Pattern
  - Frequently used in UI toolkits; example:



- P&C coupled in space and time, decoupled in flow
- Producers have to take care of subscription management and error handling


# Concepts: Message Queues

- Each message has only one consumer
- Receiver acknowledges successful processing of message
- No timing dependencies between sender and receiver
- Queue stores message (persistently), until
  - It is read by a consumer
  - The message expires (leases, TTL)



# Concepts: Publish/Subscribe

- Here: **Topic-based** Publish/Subscribe
  - Interested parties can subscribe to a topic (channel)
  - Applications post messages explicitly to specific topics
- Each message may have multiple receivers
- **Full decoupling** in space, time, and flow



# Terms

- **Event:** Any happening in the real world or any kind of state change inside an information system that is observable
- **Notification:** The reification (ger. *Verdinglichung*) of an event as a data structure
- **Message:** Transport container for notifications and control messages

# Classification (I)

- Messaging domain
  - Point-to-point (producer → consumer)
  - Subscription-based pub/sub
  - Advertisement-based pub/sub
- Subscription mechanism
  - Channel-based (=topic-based) subscription
  - Content-based subscription
  - Subject-based subscription (limited form of content-based sub.)
- Server topology
  - Single server, hierarchical, acyclic peer-to-peer, generic peer-to-peer

# Classification (II)

- Event data model
  - Untyped vs. typed vs. object-oriented
- Event filters
  - Expressiveness and flexibility of subscription language
    - Simple expressions
    - SQL-like query language
  - Evaluated in event system (router/broker network)
- **Note: scalability ↔ expressiveness tradeoff**
  - Simple expressions permit filter merging ↔ better scalability

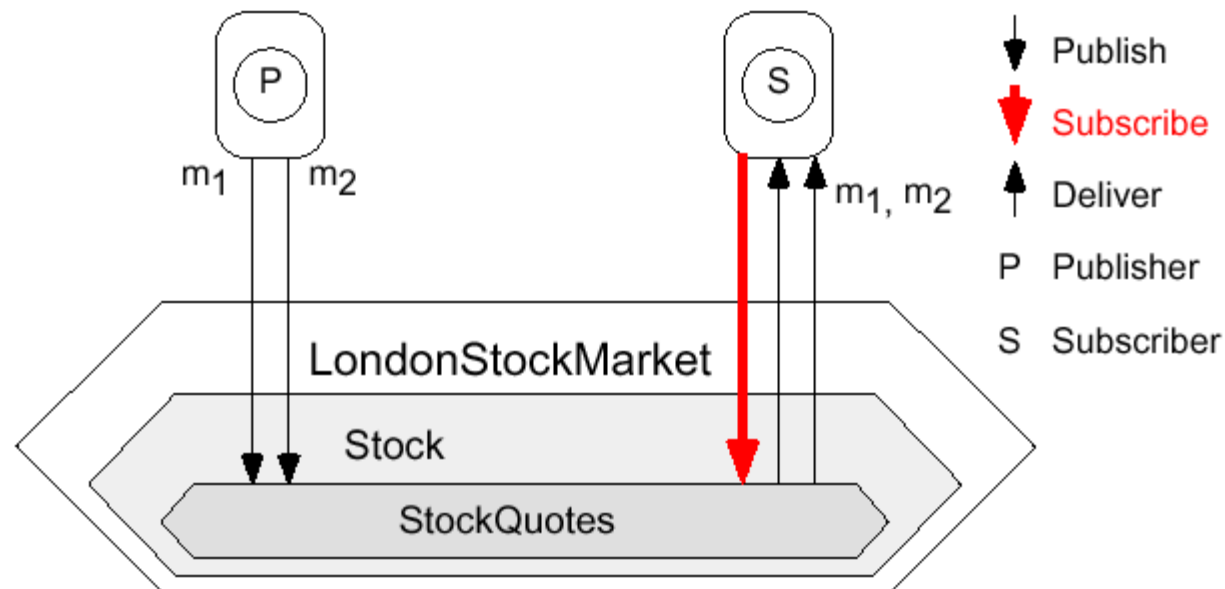
# Classification (III)

- Features
  - Scalability
  - Security
  - Client mobility
    - Transparent vs. native vs. external
  - Disconnection
  - Quality of Service (QoS)
    - Reliability & response time (real-time constraints)
  - Transactions
  - Exception handling

# Addressing

[Eugster, op. cit.]

- Channel-based addressing (=topic-based)
  - Interested parties can subscribe to a channel (analogous to mailing list)
  - Application (producer) posts messages explicitly to a specific channel
  - Channel Identifier is only part of message visible to event service
  - There is no interplay between two different channels





# Addressing (II)

- Channel-based addressing (=topic-based)
  - Extension: **topic hierarchies** (e.g., SwiftMQ)  
    <roottopic>.<subtopic>.<subsubtopic>
  - Messages are published to addressed node and all subnodes  
    iit.sales -> iit.sales.US, iit.sales.EU
  - Subscribing means receiving messages addressed to this node, all parent nodes and all sub nodes:
    - Subscription to iit.sales
    - Receives from: iit, iit.sales, iit.sales.US, iit.sales.EU
    - But not from: iit.projects
  - Subscriber receives each message only once

# Addressing (III)

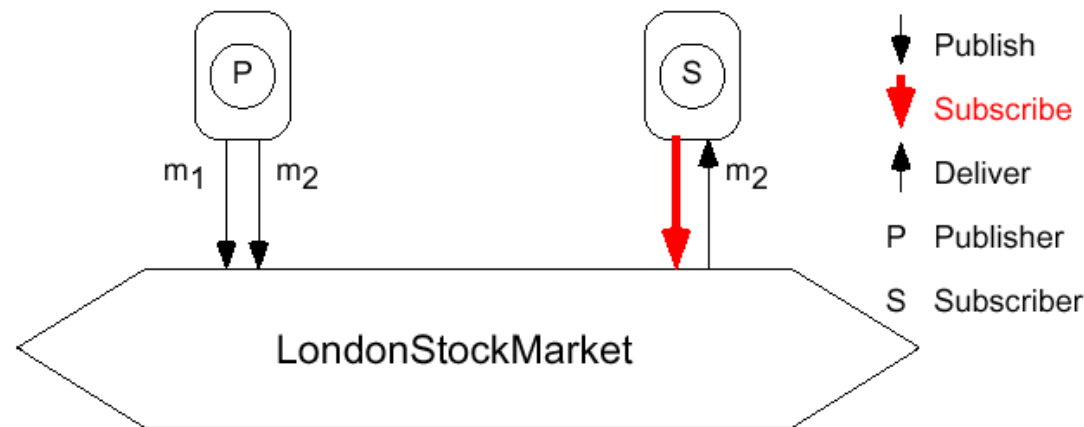
- **Subject-based addressing**
  - Limited form of content-based subscription
  - Notifications contain a **well-known attribute** – the **subject** – that determines their address
  - Subscriptions **express interest** in subjects by some form of **expressions** to be evaluated against the subject
  - Subject is
    - List of strings (e.g., TIB/Rendezvous, JEDI)
    - Properties: typed key/value-pairs (e.g., Java Messaging Service (JMS))
  - Subject (= header of notification) is visible to event service, remaining information is opaque
  - Subscription is
    - (Limited form of) regular expressions over strings (TIB, JEDI) or subset of SQL92 queries (JMS)
  - Filtering is done in the event system (router/broker network)!

# Addressing (IV)

[Eugster, op. cit.]

- **Content-based subscription**

- Domain of filters extended to the whole content of notification (i.e., payload)
- More freedom in encoding data upon which filters can be applied
- More information for event service to set up routing information



$m_1: \{ \dots, \text{company: "Telco", price: 120, } \dots, \dots \}$

$m_2: \{ \dots, \text{company: "Telco", price: 90, } \dots, \dots \}$

# Addressing (V)

- Content-based addressing

- Content-based ↔ subject-based**

- Subject-based requires some preprocessing by publisher
    - Information that might be used by subscribers for filtering must be placed in header fields
    - Thus producer makes assumptions about subscribers' interests
  - Content-based
    - Subscribers exclusively describe their interests in filter expressions

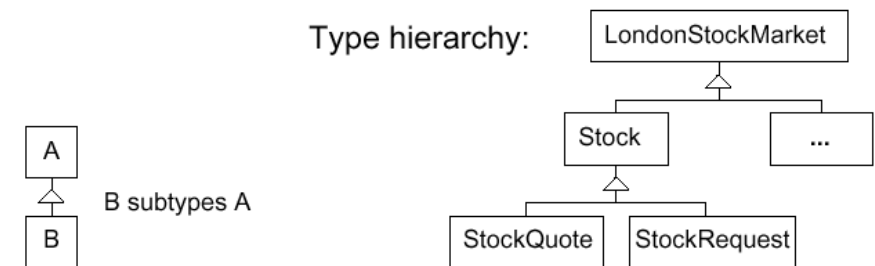
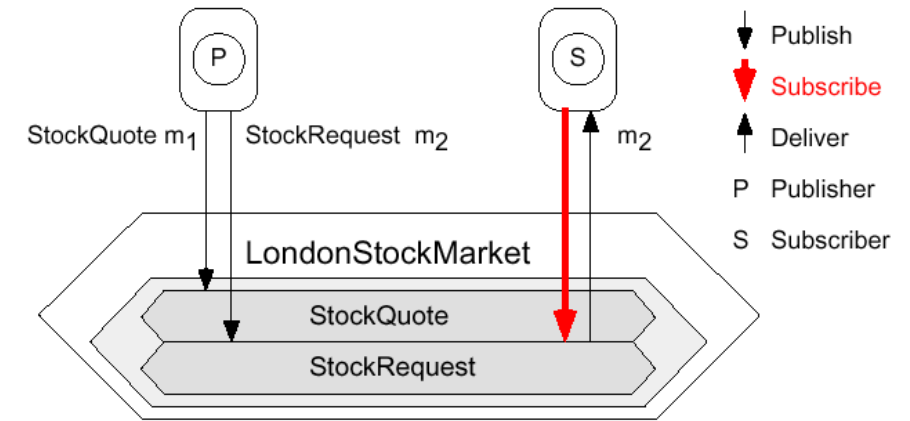
- **Concept-based** addressing

- Provides higher level of abstraction for description of subscribers' interests
  - Matching of notifications and transformation of notifications based on ontologies

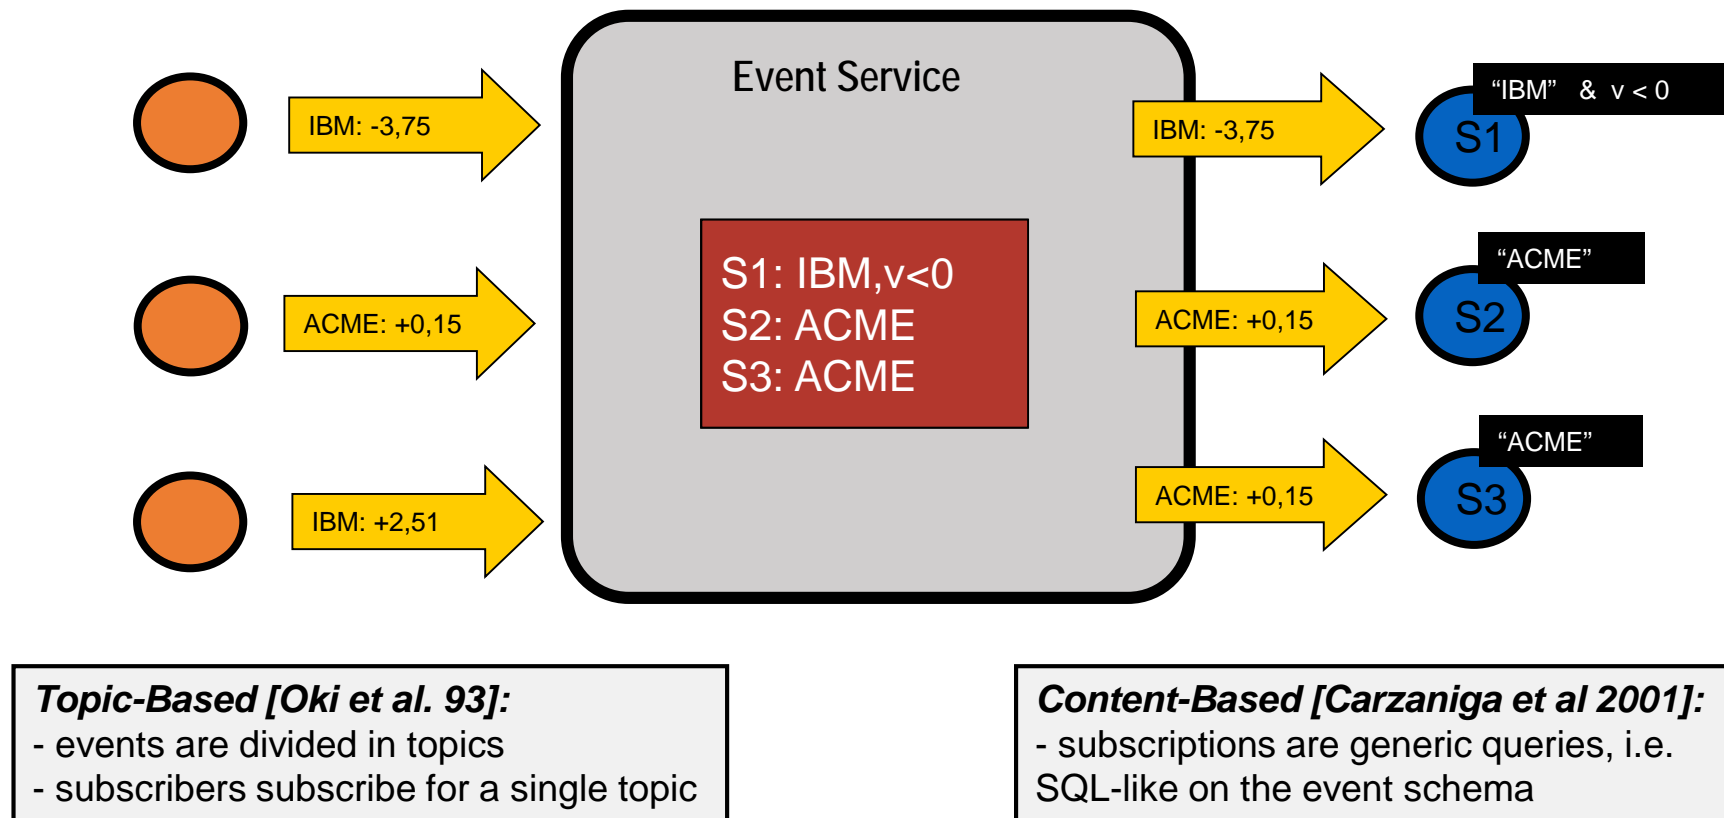
# Addressing

[Eugster, op. cit.]

- **Type-based addressing**
  - Similar to channel-based pub/sub with **hierarchies**
  - Supports subtype tests (*instanceof*)
  - Good integration of middleware & language, type safety



# Addressing - Example




# Topic vs. Content-based Subscription


- Topic-based
  - Recipients (consumer) are known a-priori
  - Many efficient implementations exist
  - Limited expressiveness
- Content-based
  - Cannot determine recipients before publication
  - More flexible
  - More general
  - Much more difficult to implement efficiently

# Subscription Mechanisms

- Subscription-based




- Advertisement-based






# Distributed Event Systems

- (Distributed) event systems
  - ... permit loosely coupled, asynchronous point-to-multipoint communication patterns
  - ... are application independent infrastructures
  - Only clients communicating via/with a **logically centralized** component



# Distributed Event Systems

- Logically centralized component
  - Single server or network of event routers (brokers)
  - Transparent (i.e. not visible) for application (=client)  
Router network can be reconfigured independently and without changes to the application → **Scalability**



# Data Model

- **Notification**

- Consists of a nonempty **set of attributes**  $\{a_1, \dots, a_n\}$
- An **attribute** is a **triple**  $a_i = (n_i, t_i, v_i)$ , where
  - $n_i$  is the attribute name
  - $t_i$  is the attribute type, and
  - $v_i$  is the value
- All data models can be mapped to this representation
  - Hierarchical messages in which attributes may be nested are flattened by using a dotted naming scheme, e.g.,
    - $\{(pos, set, \{(x, int, 1), (y, int, 2)\})\}$  can be written as  $\{(pos.x, int, 1), (pos.y, int, 2)\}$
  - Objects can be externalized (serialized) into a tree structure

# Attribute Filters

- An **attribute filter** is a simple filter that imposes a constraint on the value and type of a single attribute. It is defined as a tuple

$$A = (n, t, op, c)$$

where

- $n$  is the name of the attribute to test
  - $t$  is the expected value type,
  - $op$  is the test operator, and
  - $c$  is a constant that serve as parameter for the operator
- An attribute  $a$  **matches** an attribute filter  $A$ , iff (if and only if, gdw.)

$$a \in A :\Leftrightarrow n_A = n_a \wedge t_A = t_a \wedge op_A(v_a, c_A)$$

# Filters

- A **filter** is a stateless boolean predicate  $F(n) \rightarrow \{true, false\}$  that is applied to a notification  $n$
- If  $F(n)$  evaluates to *true*, we say that notification  $n$  matches filter  $F(n)$
- Filters that only consist of a single attribute filter are called simple filters, and filters containing multiple attribute filters are called compound filters
- **Compound filters** of the form  $F = A_1 \wedge \dots \wedge A_n$  that only contain conjunctions, are called conjunctive filters
- A notification  $n$  **matches** a filter  $F$ , iff it satisfies all attribute filters of  $F$ :  
$$n \models F :\Leftrightarrow \forall A \in F : \exists a \in n : a \models A$$
- Arbitrary logic expressions can be written as conjunctive filters in one or multiple subscriptions

# Matching: Example

## Filter

String event=alarm     ***matches***

String event=alarm     ***not matches***  
Integer level>3

## Message

String event=alarm  
Time date=02:40:03

String event=alarm  
Time date=02:40:03

# Covering (for interested reader)

- Covering between attribute filters:

- An attribute filter  $A_1$  *covers* another attribute filter  $A_2$ , iff

$$A_1 \supseteq A_2 :\Leftrightarrow n_1 = n_2 \wedge t_1 = t_2 \wedge L_A(A_1) \supseteq L_A(A_2)$$

- where  $L_A$  is the set of all values that cause an attribute filter to match

$$L_A(A_i) = \{v \mid op_i(v, c_i) = true\}$$

- Covering between filters:

- A filter  $F_1$  *covers* another filter  $F_2$ , iff for each attribute filter in  $F_1$  there exists an attribute filter in  $F_2$  that is covered by the attribute filter in  $F_1$ :

$$F_1 \supseteq F_2 :\Leftrightarrow \forall i \exists j : A_{1,i} \supseteq A_{2,j}$$

- The covering relations are required to **identify and merge similar filters**

# Overlapping (for interested reader)

- The filters  $F_1$  and  $F_2$  are **overlapping**, iff





$$F_1 \sqcap F_2 :\Leftrightarrow \\ \neg \exists A_{1,i}, A_{2,j} : (n_{1,i} = n_{1,j} \wedge (t_{1,i} \neq t_{1,j} \vee L_A(A_{1,i}) \cap L_A(A_{2,j}) = \emptyset))$$

- The overlapping relation is required to implement advertisements.
- When an advertisement A overlaps with a subscription S, we say that A **is relevant** for S.
- As a consequence, all notifications published by the client that issued A must be forwarded to the clients that issued S.



# Router Topologies

Centralized Server  
(Elvin3)



# Routing of Requests

- The network of brokers forms an **overlay network**
- Routing can be split up into two layers
  - At the lower level, requests, i.e. control messages and notifications must be routed between brokers
  - At the higher level, notifications must be routed according to subscriptions and advertisements
- Routing algorithm depends on overlay structure
  - Unstructured, generic peer-to-peer networks must avoid routing messages in cycles, e.g., use
    - Variants of distance vector routing
    - Spanning tree
  - Structured peer-to-peer networks, e.g., use
    - Distributed hash tables


# Routing: Principles

- Naive approaches:
  - **Notification flooding**
    - Flooding of notifications to all brokers in overlay network
    - Each subscription stored only in one place within the broker overlay network
    - Matching operations equal to the number of brokers
  - **Subscription flooding**
    - Each subscription stored at any place within the broker overlay network
    - Each notification matched directly at the broker where the notification enters the broker overlay network

# Routing: Principles

- **Downstream duplication**

- Route notification as a single copy as far as possible




- Clients B, C subscribe at routers 5, 6 with filter  $F_x$
- Client A publishes notification  $n_x$  (which is covered by  $F_x$ ) to router 1
- The notification is replicated not before router 4

# Routing: Principles

- **Upstream filtering**

- Apply filters upstream (as close as possible to source)



- Clients B, C subscribe at routers 5, 6 with filter  $F_x$
- Client A publishes notification  $n_y$  (not covered by  $F_x$ ) to router 1
- The notification is discarded at router 1

# Routing with Subscriptions

- Each broker maintains a **routing table**  $T_s$  to route notifications based on subscriptions
- Routing of notifications: A notification  $n$  is only forwarded to a destination  $D$ , iff

$$\exists (D, F) \in T_s : n \in F$$

- Routing of subscriptions: If a subscribe or unsubscribe request is received, the table  $T_s$  is updated accordingly.
  - Subscribe or unsubscribe requests are potentially forwarded to all neighbors according to the underlying routing algorithm

# Routing with Advertisements

- **Basic Idea**

- Subscriptions are only forwarded towards publishers that intend to generate notifications that are **potentially relevant** to this subscription
- Every advertisement is forwarded throughout the network, thereby forming a tree that reaches every server
- Subscriptions are propagated in reverse, along the path to the advertiser, thereby **activating** the path
- Notifications are then forwarded only through **activated** paths.

# Scalability

- System should be **scalable** in terms of
  - the number of clients (i.e., producers and consumers),
  - the number of event routers (brokers),
  - the number of subscriptions and advertisements, and
  - the amount of traffic (e.g., number of notifications/second)
- Problems in unstructured peer-to-peer overlays
  - Either subscriptions or advertisements forwarded to each node
    - Assumption (for Internet-based services): advertisements are rather static, subscriptions are dynamic → use routing with advertisements
  - Routing tables grow proportionally with the size of the network
    - use filter merging
    - use structured overlays




# Filter Merging (for interested reader)

- Inexact Merging

$F_M$  is an inexact merge of  $F_1$  and  $F_2$  iff


$$F_M \sqsupseteq F_1 \wedge F_M \sqsupseteq F_2$$



- Exact Merging

$F_M$  is an exact merge of  $F_1$  and  $F_2$  iff

$$F_M \sqsupseteq F_1 \wedge F_M \sqsupseteq F_2 \wedge \neg \exists F_3 : (F_3 \not\sqsupseteq F_1 \wedge F_3 \not\sqsupseteq F_2 \wedge F_M \sqsupseteq F_3)$$



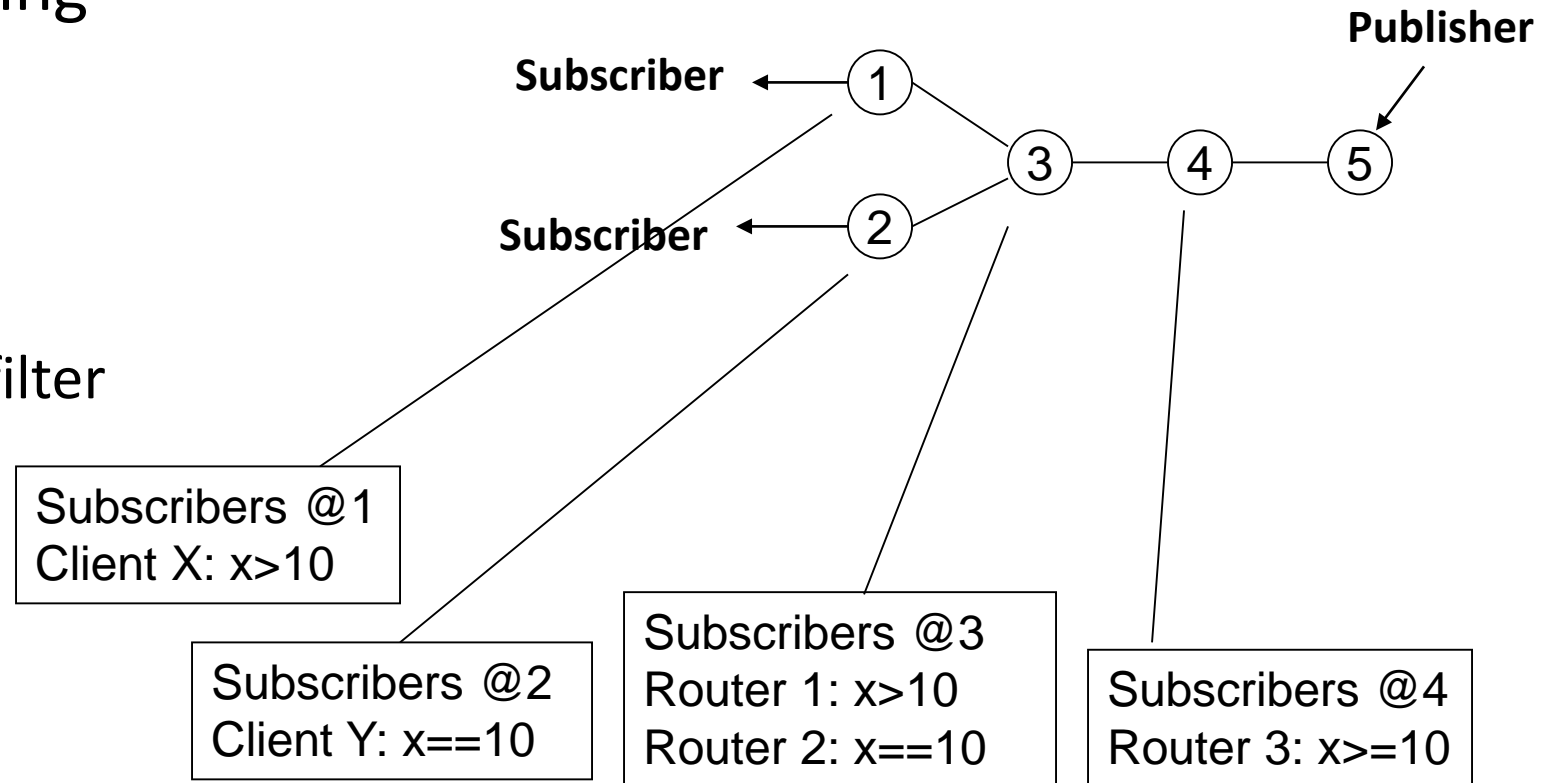
# Filter Merging: Example

- Filter merging

- Filter X  
 $x > 10$

- Filter Y  
 $x == 10$

- Merged filter  
 $x \geq 10$



# Structured Overlays

- Systems based on **distributed hash tables** (e.g. SCRIBE)
- In a DHT, the storage location of an information item is defined by its hash value
  - Channel-based addressing: calculate hash value from channel name
  - Content-based addressing: no general solution
    - „Channelization“: calculate hash from selected attributes, e.g. message type
- The (global) subscription table is distributed over the network
  - A broker is responsible for specific subscriptions
  - The broker is the rendezvous point for publishers and subscribers

# Structured Overlays (II)

- Routing of subscriptions
  - Subscriber calculates hash of subscription  $h(S)$  and sends it to the broker with hash  $h(B)$  closest to  $h(S)$ . The subscription is stored at B.
- Routing of notifications
  - Publisher calculates hash of notification  $h(n)$  and sends it to the broker with  $h(B)$  closest to  $h(n)$ . Broker B has a list of all relevant subscribers.

# QoS and Transactions


- Quality of Service

- Guaranteed delivery
  - Logistics
  - Stock quotes
- Low latency
  - Sensor, audio, or video data streams

- Local Transactions

- between the publisher and the event service, or between the event service and the subscriber
- groups a series of operations into an atomic unit of work

# Mobility Support: Durable Subscriptions



- Messages are stored for each subscriber
  - Permits disconnection of subscriber
  - **But:** subscriber bombarded with messages on reconnect (remedy: use TTL)

# Generic Pub/Sub Interface

- **subscribe():**
  - register interest in events to event service  
(without specifying the effective sources of these events)
  - usually not forwarded to publishers
- **unsubscribe():**
  - terminates a subscription
- **publish():**
  - generates an event at the event service
  - usually propagates event to all relevant subscribers
- **advertise():**
  - expresses the intent to publish a particular kind of event
  - can be used for discovery of publishers or to simplify matching
  - not strictly needed and often not implemented

# System Examples

- Industry-strength
  - **JMS**
  - CORBA Notification Service
  - **ZeroMQ**
  - Elvin
  - IBM WebSphere MQ Event Broker (Gyphon)
- Academic prototypes
  - REBECA
  - SIENA



# JMS: Java Message Service

- API - „Common set of interfaces and associated semantics“
- Domains
  - Point-to-point: message-queue
  - Publish/subscribe
    - Topic-based
    - Subject-based
    - Durable subscribers
- Separated administration
  - Queues and topics are created with product-specific administration tools
  - Application independent
  - Support of local transactions

# JMS: Java Message Service

- Message format
  - **Header**: predefined fields (ID, destination, timestamp, priority)
  - **Properties** (optional): accessible for filtering  
values can be boolean, byte, int, ... double and string
  - **Body** (optional): five types
    - TextMessage: String (XML Document)
    - MapMessage: Key/Value-Pairs
    - BytesMessage: Stream of uninterpreted bytes
    - StreamMessage: Stream of primitive values
    - ObjectMessage: A serializable object
- Event consumption
  - Synchronously: subscriber explicitly fetches message from destination
  - Asynchronously: subscriber registers a message listener

# JMS: Message Filtering

- SQL92 conditional expressions (limited)
  - Logical operators in precedence order: NOT, AND, OR
  - Comparison operators: =, >, >=, <, <=, <> (not equal)
  - Arithmetic operators in precedence order: +, - (unary) \*, / (multiplication and division) +, - (addition and subtraction)
  - arithmetic-expr1 [NOT] BETWEEN arithmetic-expr2 AND arithmetic-expr3 (comparison operator)
  - identifier [NOT] IN (string-literal1, string-literal2,...) (comparison operator where identifier)
  - identifier [NOT] LIKE pattern-value [ESCAPE escape-character]
  - identifier IS [NOT] NULL (comparison operator that tests for a null header field value or a missing property value)

# JMS: Message Filtering (II)


- Examples:
  - NewsType='Opinion' OR NewsType='Sports'
  - phone LIKE '12%3'
  - JMSType='car' AND color='blue' AND weight>2500

# JMS: SwiftMQ

- Domain
  - Point-to-point
  - Topic- and subject-based publish/subscribe
- Server topology
  - Generic peer-to-peer: federated router network
- Features
  - Fully implements JMS 1.0.2 specification
  - Topic hierarchies
  - SQL-Like predicate topic addressing  
Permits subscription with topic name wildcard. Example:  
iit.s%s.\_S matches iit.sales.US
  - File based persistent message store

# SIENA

- SIENA = Scalable Internet Event Notification Architecture
- Domain
  - Advertisement-based publish/subscribe
  - Content-based subscriptions
- Server topology
  - Generic peer-to-peer
  - Hybrid topology
    - LAN: hierarchical
    - WAN: generic peer-to-peer
- Data model
  - Notification is set of attribute=(name, type, value)
  - Limited set of types (string, time, date, integer, float, ...)
- Subscription language
  - Filter is set of attr\_filter=(name, type, operator, value)
  - Operators: any, =, <, >, >\* (prefix), \*< (postfix)




# PADRES

- PADRES
  - Uses **advertisement based routing** which is adopted from SIENA and REBECA system.
  - Consists of a broker network.
  - Publishers and subscribers are connected to brokers as clients.
  - A publisher advertises before it publishes. The advertisement is flooded in the broker network.

# PADRES (II)

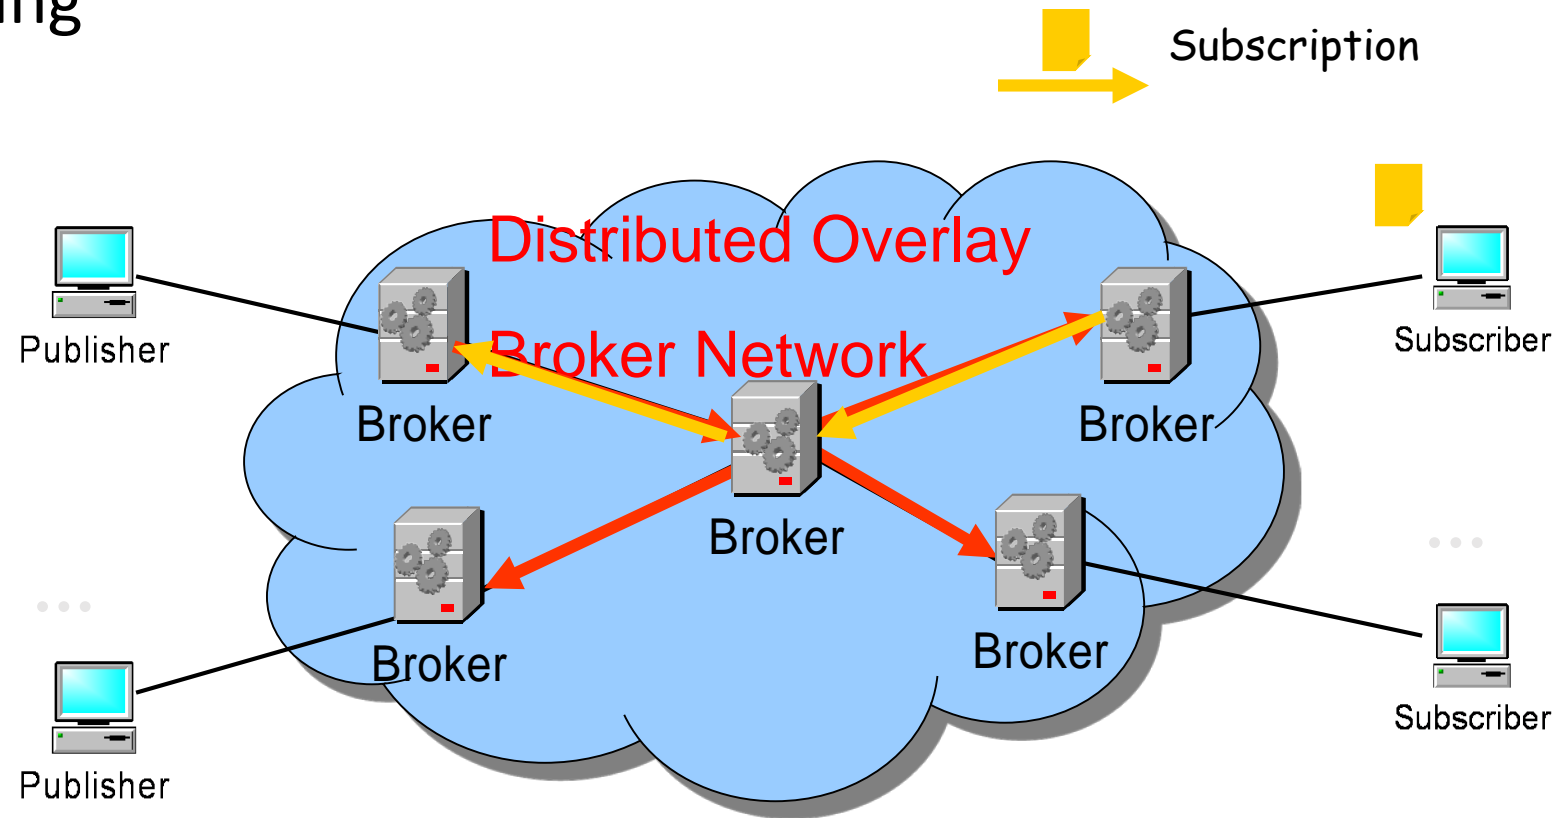
- Advertising






# PADRES (II)

- Subscribing



# PADRES (II)

- Publishing



# Elvin

- Elvin3
  - Subscription-based publish/subscribe
  - Content-based subscriptions
  - Centralized server
- Elvin4
  - Data model: typed key/value-pairs  
Types: integer (32 and 64), string, FP, binary data (opaque)
  - Subscription Language:  
Simple Integer and FP arithmetic  
Strings: POSIX ERE (Extended Regular Expressions),  
begins-with, ends-with, contains (for better optimization)
  - Source code available for non-commercial use
  - Proxy at network boundary to support disconnection

# ZeroMQ API

- **An attempt to create a unique interface for different interaction models**
- Open source (<http://zeromq.org/> )
- ZeroMQ socket: multiple types for different patterns:
  - Request/Response, Publish/Subscribe, etc.
- Uniform API across different programming languages:
  - More than 40 bindings (including C, Erlang, Go, Java, Python, etc.)
- Addressing similar to URL: string [transport]://[address]
  - Automatic DNS resolving

# ZeroMQ API (II)

- ZeroMQ offers asynchronous messaging with transparent message queuing:
  - Sending messages does usually not block
  - Messages are queued at sender if destination host is not available or too slow until “high watermark” is reached (queue full), then blocks or discards depending on socket type
  - Can reconnect automatically (e.g. if server temporarily unavailable) when underlying transport signals closed connection

# Berkeley Socket API vs. ZeroMQ API

Berkeley Sockets	ZeroMQ equivalent
Socket	zmq_socket()
Bind	zmq_bind()
Listen	-
Accept	-
Connect	zmq_connect()
Send	zmq_send()
Receive	zmq_recv()
Close	zmq_close()

- Most API calls map to Berkeley sockets, except:
  - Sequence Bind/Listen/Accept from Berkeley sockets passive side replaced by the ZeroMQ bind

# Publish/subscribe with ZeroMQ

- **Publisher** and **subscriber** use **different sockets types**:
  - PUB for publishers, SUB for subscribers
- Topic based filtering: message prefix matching
- Routing: subscription flooding to publishers (filtering at publishers)
- Topology: PUB and SUB sockets can be connected directly or through brokers (which use XPUB/XSUB sockets and act like a single subscriber or publisher to connecting sockets)
- Subscriptions are aggregated at brokers and only forwarded once

# Publish/subscribe Interface using ZeroMQ

General pub/sub API	ZeroMQ equivalent
Socket	zmq_socket()
Bind	zmq_bind()
Connect	zmq_connect()
Publish	zmq_send()
Subscribe	zmq_setsockopt(ZMQ_SUBSCRIBE)
Unsubscribe	zmq_setsockopt(ZMQ_UNSUBSCRIBE)
Close	zmq_close()



# Example of pub/sub with ZeroMQ's API

- Usage of C binding:

```
void *context= zmq_ctx_new ();  
void *publisher = zmq_socket(context, ZMQ_PUB);  
zmq_bind(publisher, "tcp://localhost:5001");  
  
zmsg_send("Hello World", publisher);  
  
zmq_close(publisher );  
zmq_ctx_destroy(context);
```

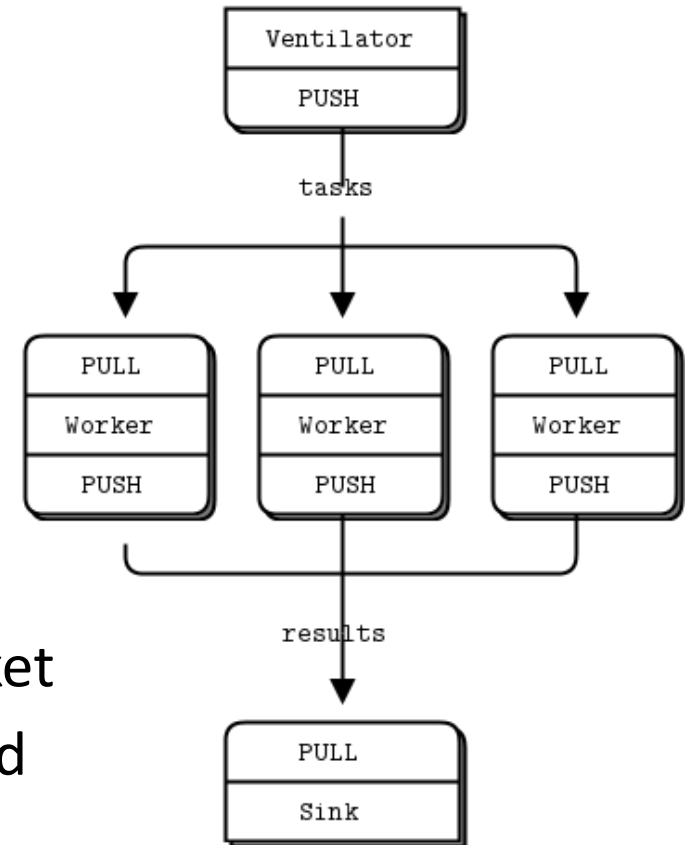
← Publisher

Subscriber →

```
void *context = zmq_ctx_new ();  
void *subscriber = zmq_socket(context, ZMQ_SUB);  
zmq_connect(subscriber, "tcp://localhost:5001");  
  
zmq_setsockopt (socket, ZMQ_SUBSCRIBE, "H", 0);  
zmsg_t *msg = zmsg_rcv(subscriber);  
  
zmq_close(subscriber);  
zmq_ctx_destroy(context);
```

# ZeroMQ pipeline messaging pattern

- Intended for **task distribution**
- Tasks are distributed from a ventilator to workers and are aggregated at a sink
- Ventilator uses PUSH socket
  - Sending messages round-robin to workers
- Workers use PULL socket to get tasks
  - Blocking until task is available
  - Queues incoming tasks if busy
- Workers use PUSH Socket to send results to sink's PULL socket
- Because of rigid round-robin distribution only recommended for task that execute quickly



# Payload Serialization Formats

- How can we convert complex data objects into a sequence of bits, to be used as part of service request or response messages?
  - Use of “Data Serialization / Encoding Formats”
- Conflicting requirements for data serialization formats
  - Efficient encoding and decoding (time and space/size)
    - Leads to binary encodings
  - Extensibility and backward-compatibility
    - Enables decoupling of the server and client implementations
  - Human readability
    - Simplifies development by increased visibility
  - Type safety
    - Automatic detection of type errors
  - Self-describing schema

# Examples of Data Serialization Formats

- JSON

- Attribute-value pairs
- Supported types: Number, String, Boolean, Array, Object, null
- Everything encoded as strings, **not very space efficient**
- Can be additionally compressed (cJSON, bJSON)
- **Human readable**, programming language independent
- Not self describing, no type checking, no schema

- Google Protocol Buffers

- **Very efficient binary** data serialization format
  - Used as payload serialization for RPC calls and as data persistence format
- **Strongly typed**, separate schema file (.proto)
- Requires a compilation step to read and serialize messages
- **Extensible and backward-compatible**

# Examples of Data Serialization Formats (II)

- Apache Avro
  - Efficient binary data serialization format
  - Simple integration in dynamic languages, no compilation needed
  - **Dynamic typing**
  - **Self-contained schema** specification (using JSON)

# Summary - Pub/Sub

- **Loosely coupled systems**

- Space decoupling
- Time decoupling
- Control flow decoupling

- **Publish/Subscribe**

- Powerful and scalable abstraction for decoupled interaction
- Problems are at the algorithm & implementation level
- Research challenges: scalability/expressiveness-tradeoff, fault tolerance, integration with P2P, security, reliability, ...